



On-device latency comparison for human action recognition

(a) Embedded end-to-end action-recognition pipelines



(b) Model/core inference only (not end-to-end)



Protocols, datasets, precision, input windows, and Jetson variants differ. E2E includes preprocessing/post-processing where reported; core/model values exclude the surrounding perception pipeline.